Machine Learning with Opponents

Brendan Herger
brendan@thisismetis.com
Slides: https://goo.gl/wY8iFY


MAN, I SUCK AT THIS GAME. CAN YOU GIVE ME A FEW POINTERS?
0x3A28213A
0x6339392C,
0x7363682E.
I HATE YOU.

Overview

Sampling
Feature Engineering
Modeling

Sampling

Observation Weighting

- Effect cost function by weighting every row at train time
- Weights include
  - Uniform weight
  - Observation age (staleness)
  - Random down-sampling
  - Up-sampling known opponent attacks

Observation Weighting

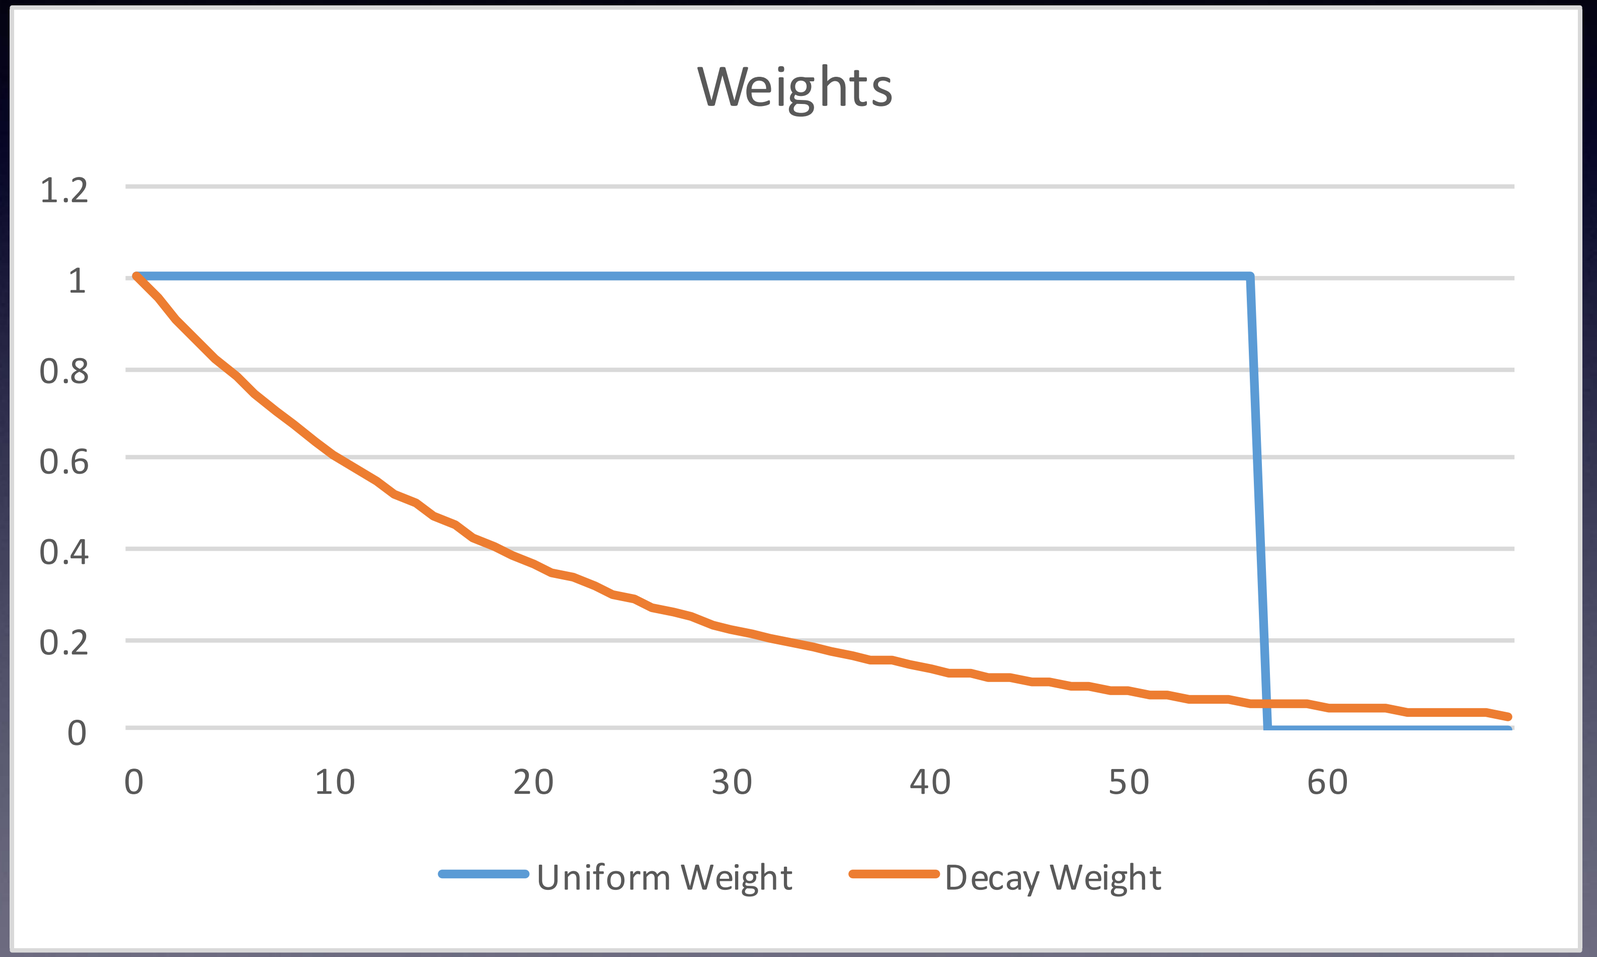

SMOTE

(Synthetic Minority Over-sampling Technique)

- Goal: Better model rare events (opponent attacks)
- Majority class: Down sample, with some probability
- Minority class: Create ‘synthetic’ observations

https://www.jair.org/media/953/live-953-2037-jair.pdf

SMOTE

1. Select minority point
2. Select neighbor
3. Create new point

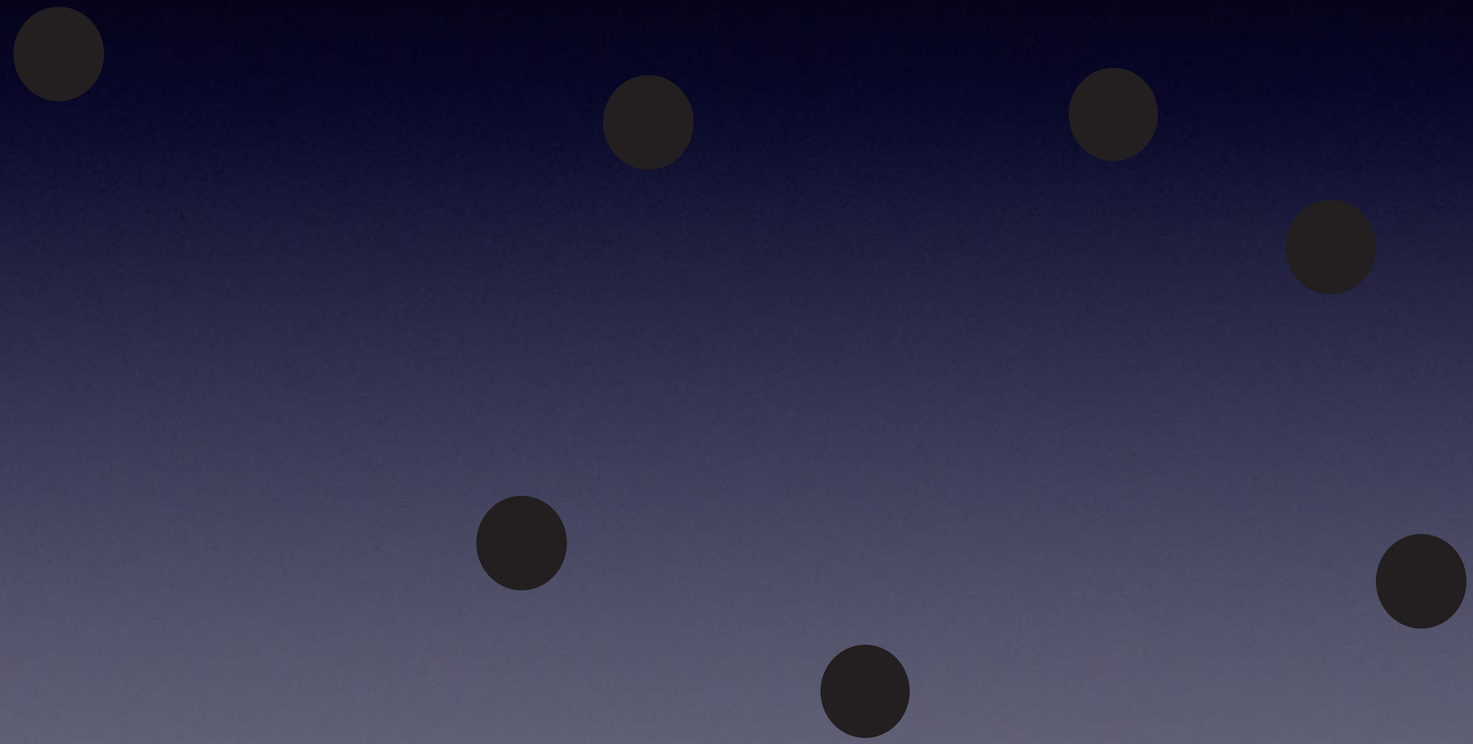

https://www.jair.org/media/953/live-953-2037-jair.pdf

SMOTE

1. Select minority point
2. Select neighbor
3. Create new point

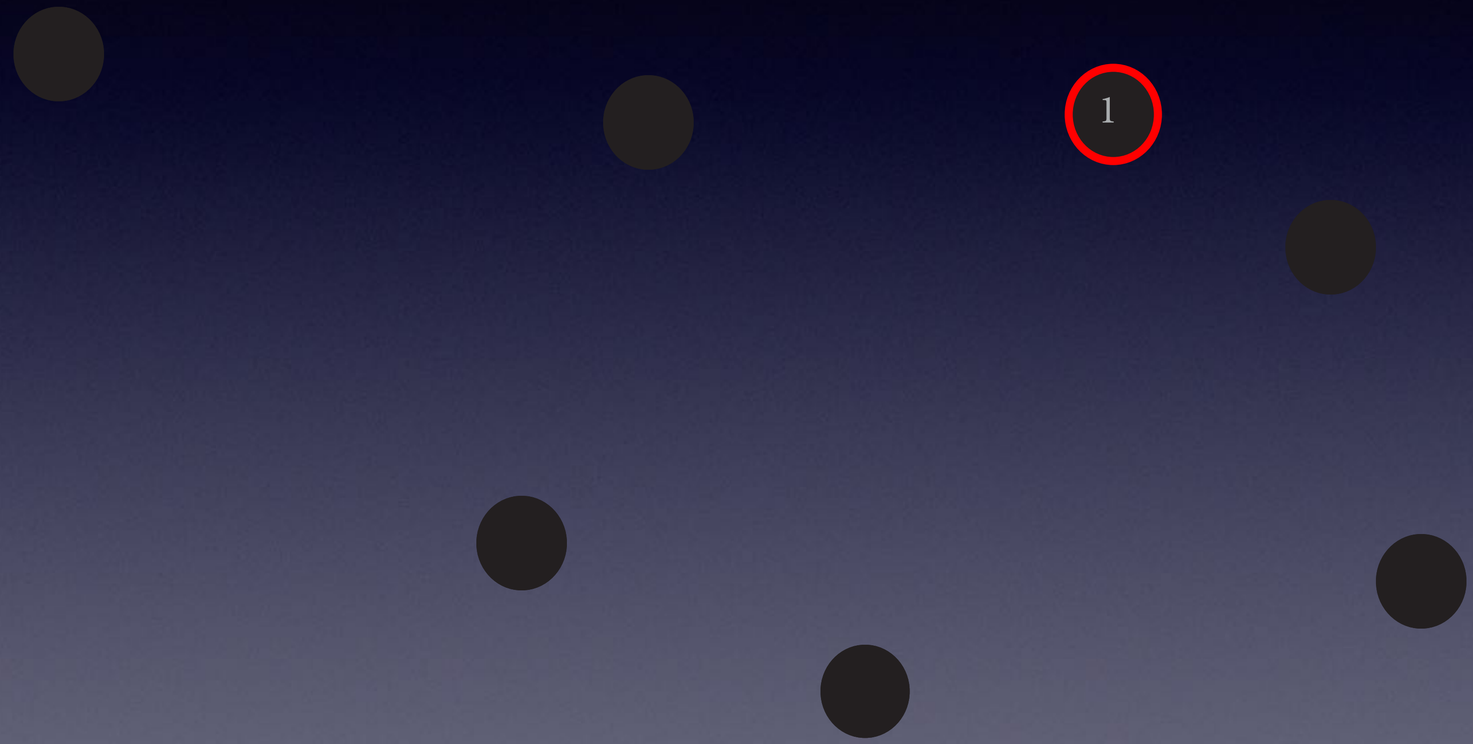

https://www.jair.org/media/953/live-953-2037-jair.pdf

SMOTE

1. Select minority point
2. Select neighbor
3. Create new point

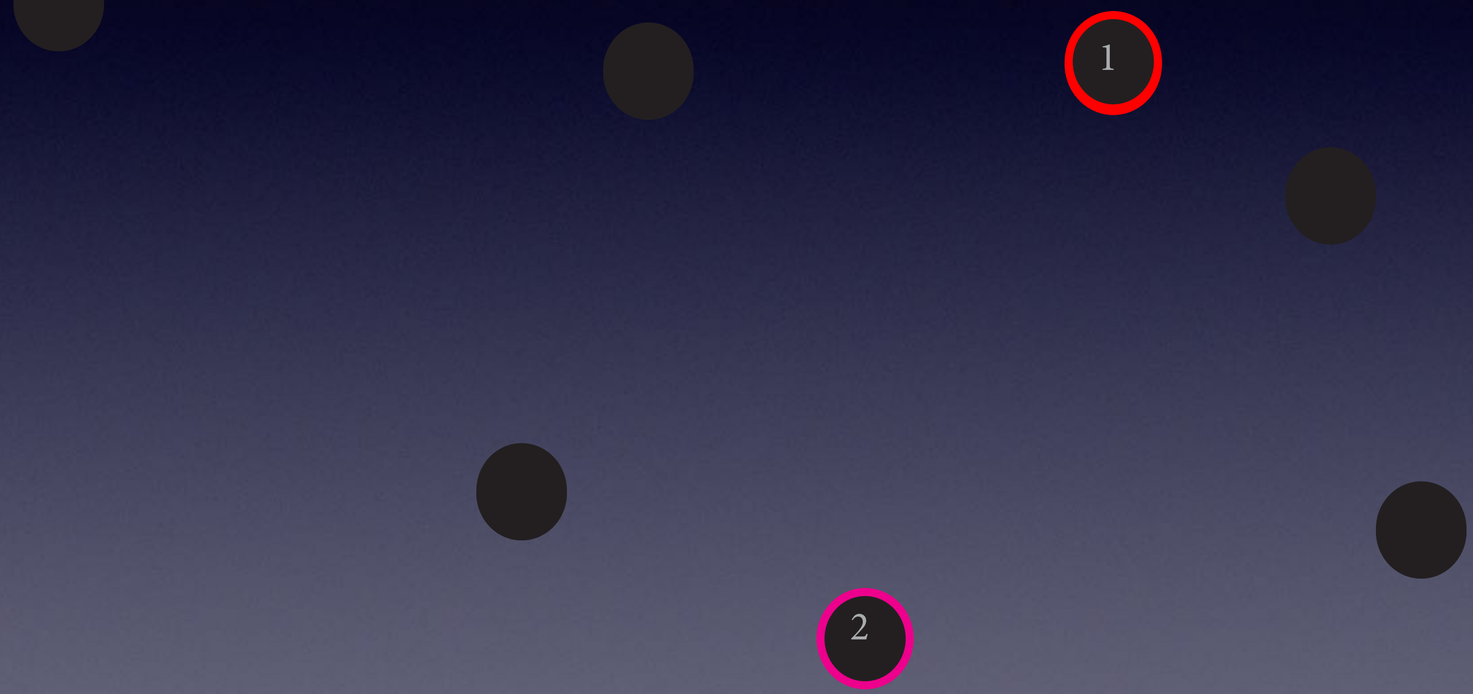

https://www.jair.org/media/953/live-953-2037-jair.pdf

SMOTE

1. Select minority point
2. Select neighbor
3. Create new point

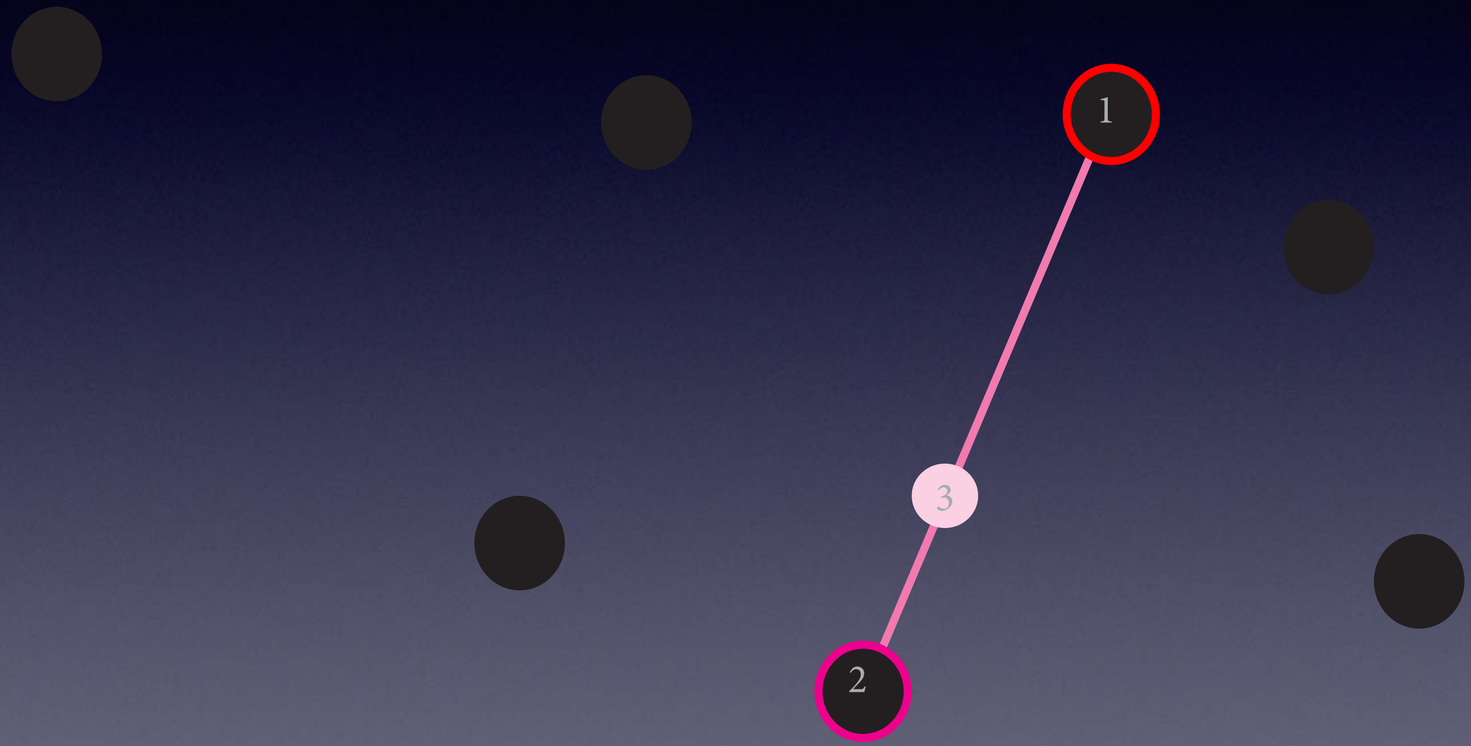

https://www.jair.org/media/953/live-953-2037-jair.pdf

Observation Weighting
SMOTE Sampling

Features

Outlier Detection

- Goal: Create outlier score
- Train learner to re-create input vector
  - PCA: Reduce dimensionality, increase dimensionality
  - Neural Network: Train auto-encoder, with bottleneck
- Measure distance from output vector to input vector

https://github.com/h2oai/h2o-training-book/blob/master/hands-on_training/anomaly_detection.md

Low Rank Models

- Goal: Reduce dimensionality for dataset with many variables
- Reduce dimensionality with generalized PCA
- Model directly on components (latent factors)

https://web.stanford.edu/~boyd/papers/pdf/glrm.pdf
https://github.com/h2oai/h2o-tutorials/blob/master/tutorials/glrm/glrm-tutorial.md

Low Rank Models

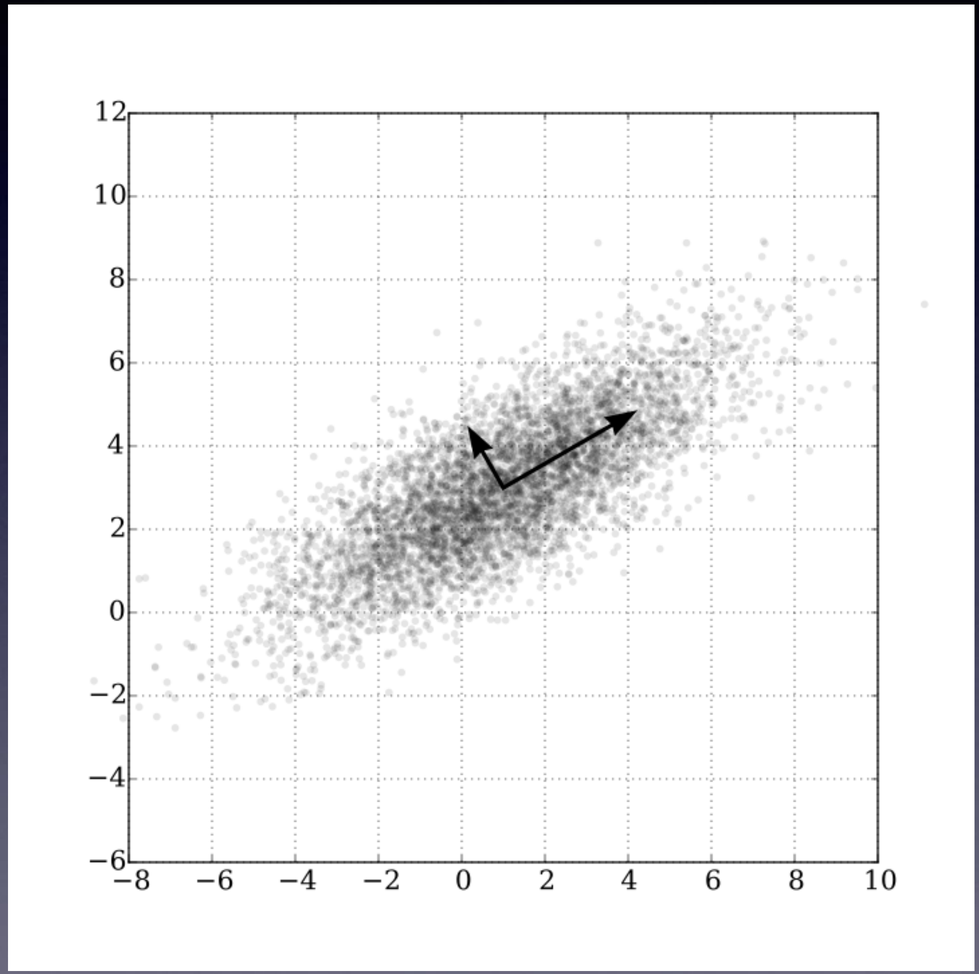

https://web.stanford.edu/~boyd/papers/pdf/glrm.pdf
https://github.com/h2oai/h2o-tutorials/blob/master/tutorials/glrm/glrm-tutorial.md

Outlier Detection

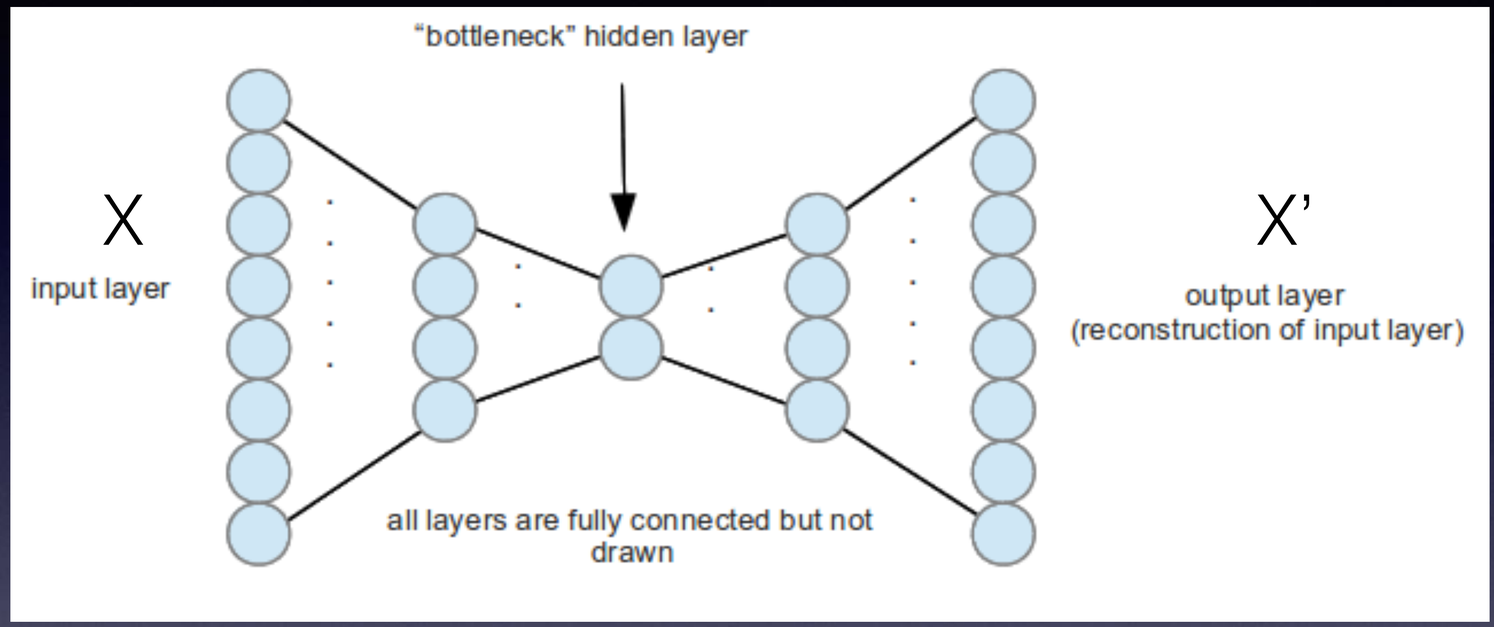

Outlier score: |X'-X|

https://github.com/h2oai/h2o-training-book/blob/master/hands-on_training/anomaly_detection.md

Outlier Detection
GLRM

Modeling

Grid search

- Goal: Find optimal hyper-parameters for given class of models
- Create every possible permutation of hyper-parameters, and compute models until heat death of universe

https://github.com/h2oai/h2o-3/blob/master/h2o-docs/src/product/tutorials/GridSearch.md

Neural Networks

- Too complicated to cover here 😞
- Strengths: Able to capture complicated, non-linear relationships. Deals well with class imbalance, near-arbitrary data types
- Weaknesses: Difficult to train, many free parameters

Ensemble Modeling

- Goal: Leverage a diverse set of algorithms
- Train multiple classes of algorithms (tree based, linear, neural network), possibly with multiple hyper-parameters, combine scores with meta model

https://github.com/h2oai/h2o-3/blob/master/h2o-docs/src/product/tutorials/GridSearch.md

Ensemble Modeling

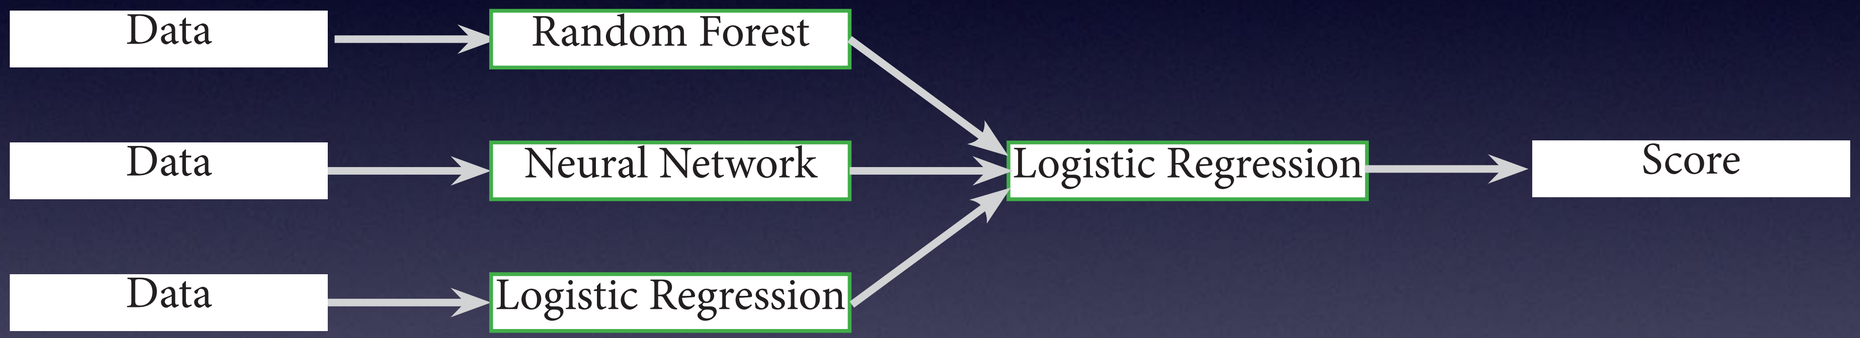

https://github.com/h2oai/h2o-3/blob/master/h2o-docs/src/product/tutorials/GridSearch.md

Grid Search
Neural Networks
Ensemble models

Thanks!

Slides: https://goo.gl/wY8iFY

brendan@thisismetis.com